



# Kernel Tales

Testing security of Kernels for Android or  
embedded devices by fuzzing drivers  
(on aarch64)



Vito Rallo  
@vitorallo

vito.rallo@gmail.com

I work for:   
**pwc**



# Agenda

- Introduction to vulnerability research in the kernel; why?
- Analysis of Android embedded devices (supply chain);
- Security considerations for Android and IoT devices in general;
- Threat modelling;
- Methodology;
- Trinity, or better to say: Dronity!
- Practical details;
- Setup your toolchain;
- Navigate and understand the source code;
- Little hacks, cross-compile, deploy, enjoy!
- Compile and run dronity... fuzz fuzz fuzz
- Fuzzing on the ODROID BOARD
- Fuzzing on the Nexus 9
- The aarch64 emulator (QEMU/ranchu)
- Fuzzing ranchu;
- Customize Dronity, develop your own IOCTL module;
- (time permitting) Debugging the kernel on aarch64 emulator: attach gdb, use unstripped symbols kernel, breakpoints and kernel panic analysis

THEORY

HANDS-ON



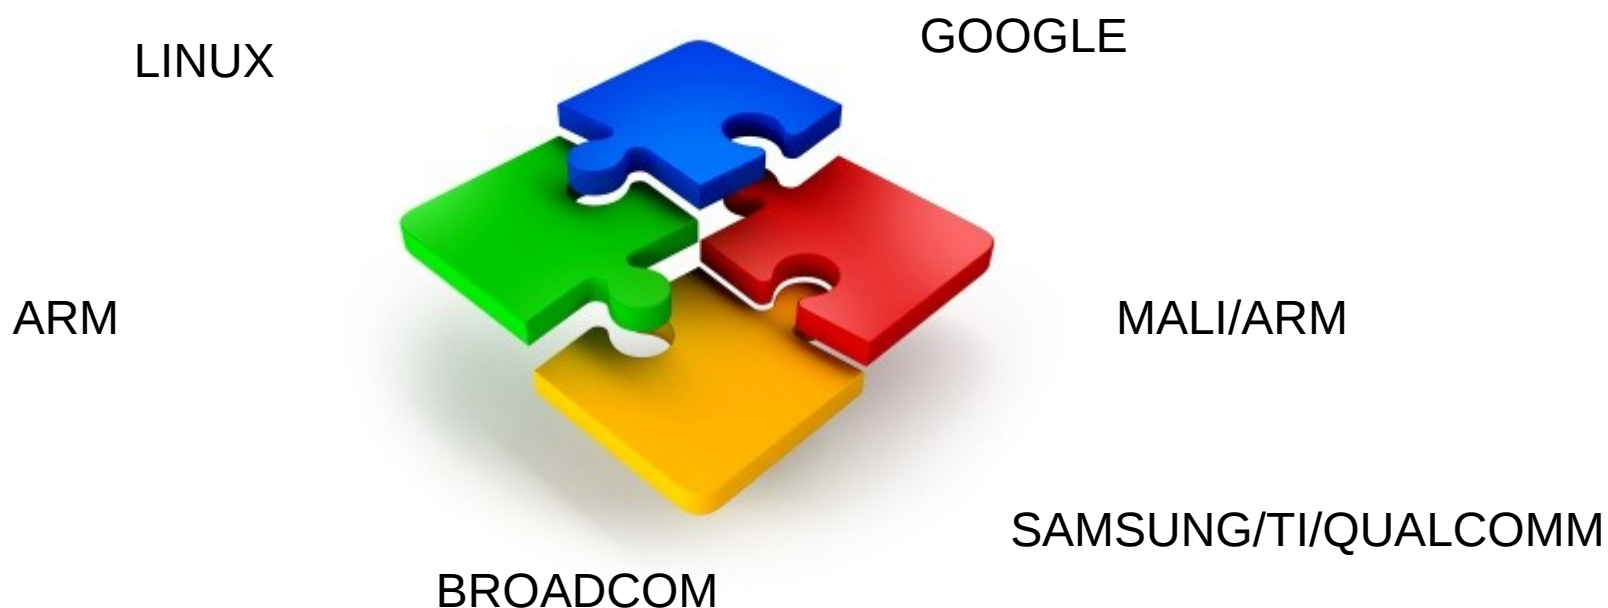
# Few more bits

- This workshop is not:
  - Kernel/bin reverse engineering
  - Exploitation!!!
- I am not:
  - A serious kernel developer
  - Full time vulnerability researcher
  - Dan Rosenberg, Geohot.. etc :)



# Puzzle Code

- A puzzle of different source, providers, vendors, where proper code review and control is hard!



## *Security in the embedded devices supply chain*



Security must be guaranteed and built in from the bottom of the embedded device **supply chain**, to create and deliver good, safe and high quality devices.

Two important factors:

- the inner security and the features built into the product;
- the way developers use them.



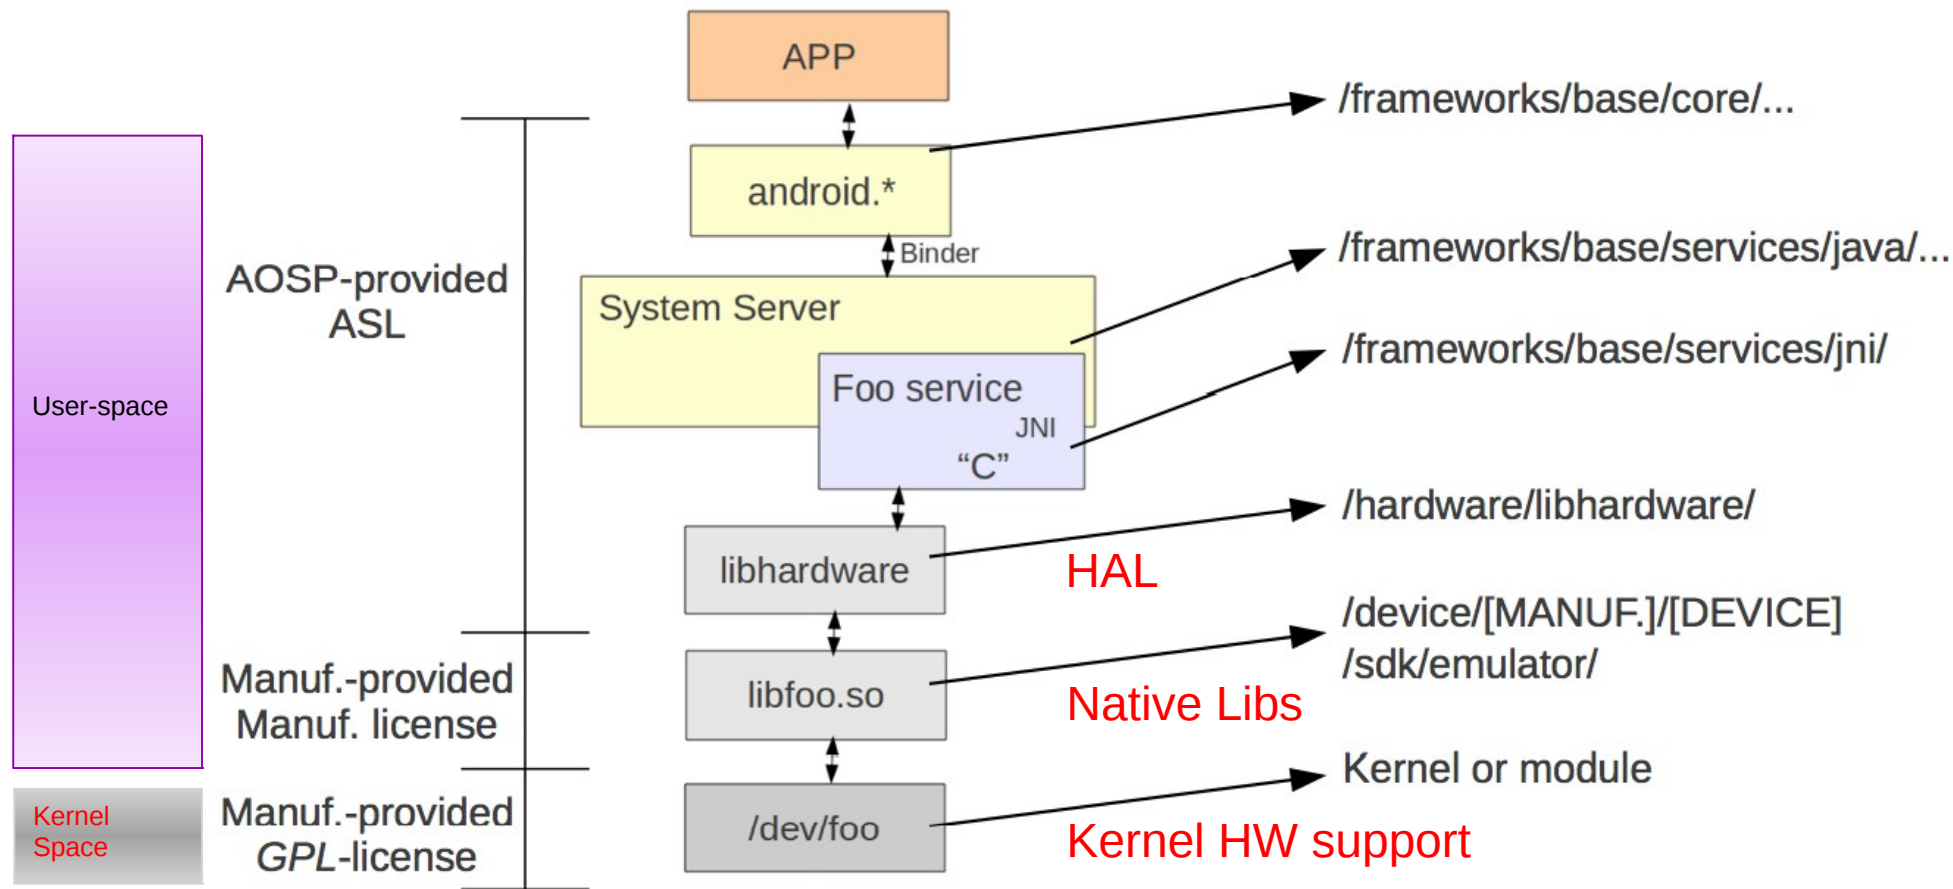


## Impact of ROOTed Android

- Customize the ROM (install/remove)
- Break DRM, use unhallowed media sources
- Connect external devices
- Get malicious software, botnets and others
- Dump proprietary code and reverse
- Crack and use software decoding tools for smartcards and encrypted channels

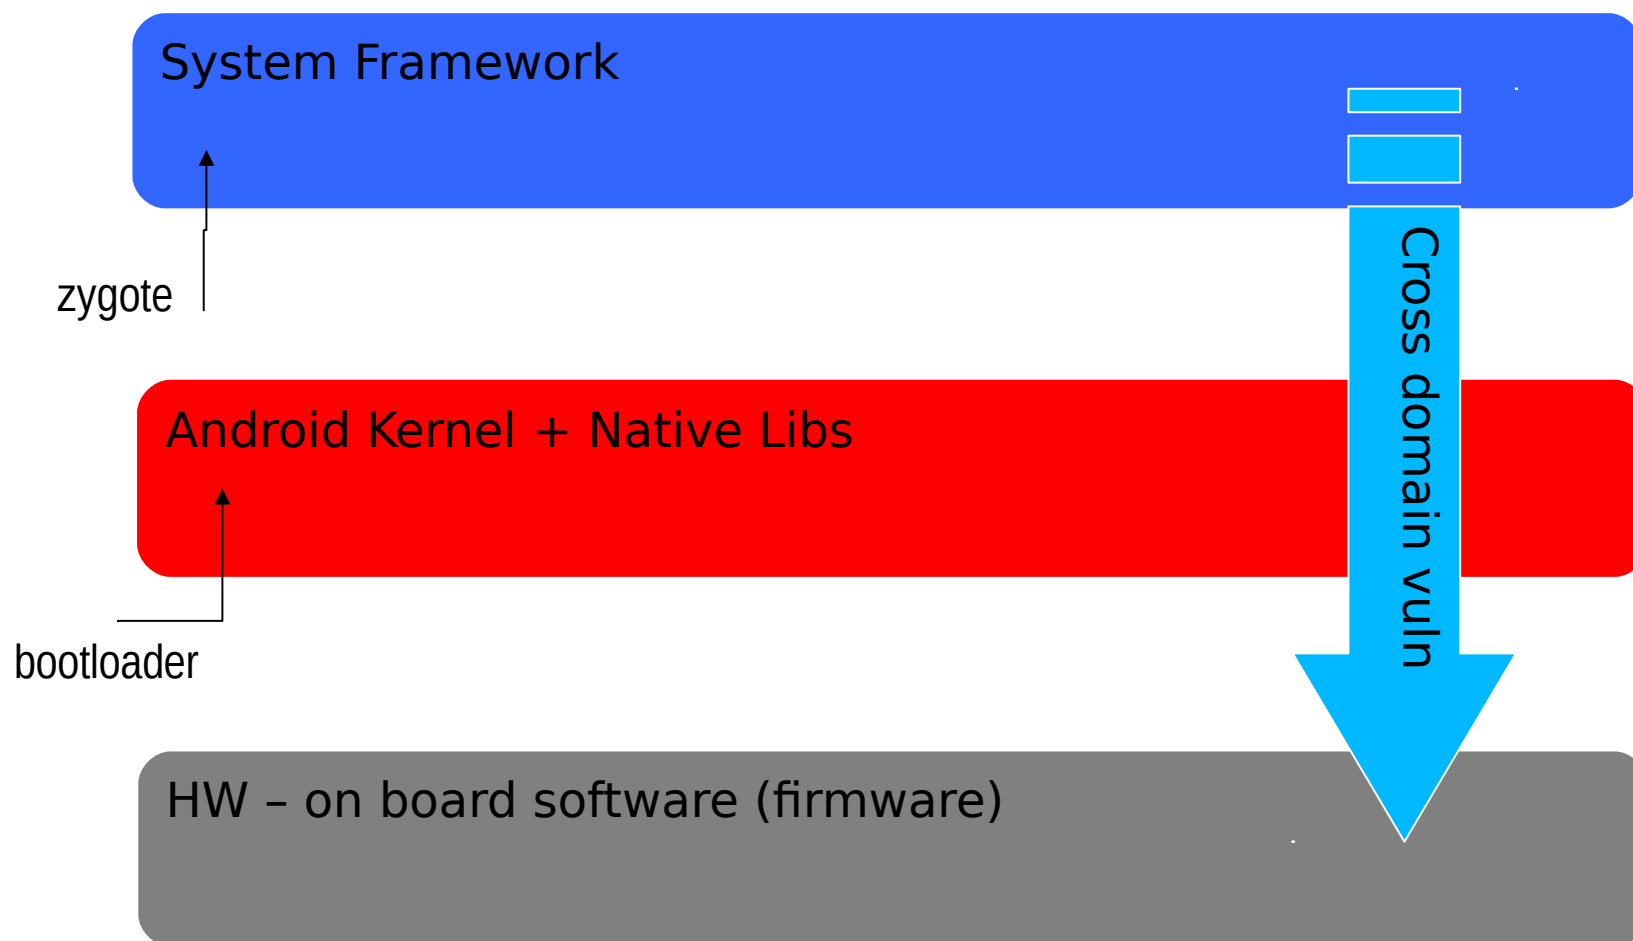


# Where are we?





# Embedded Device Security Domains

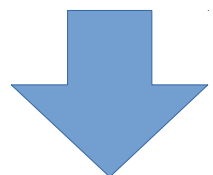




# Rooting trend on Android



- Userspace processes exploitation
- Wrong hardening and service configuration or init.script and symlinks
- Userspace libs (e.g. libstagefright)
- Recovery and update scripts
- Dalvik privilege elevation, up to system (master key family etc)



- Daemon rooting, setuid, memory corruption (gingerbread, zergrush) and other old techniques
- **Kernel root exploitation:** framaroot (dev/exynosmem), towelroot and finally the all new pingpong root CVE-2015-3636

Kernel improvements, SELinux, Samsung root mitigation, write protection (PXN), SDLC and may be one day KASLR..... **is the same attention reserved for embedded devices?**



# .. and today?

Latest bugs:

*ZipBug9950697*

*Zip Bug 8219321 / Master keys*

*Zip Bug 9695860*

*Jar Bug 13678484 / Android FakeID*

*CVE 2013-6282 / put/get\_user*

*CVE\_2011\_1149 / PSNueter / Ashmem Exploit*

*CVE\_2014\_3153 / Futex bug / Towelroot*

*CVE 2014-3847 / WeakSauce*

*StumpRoot*

*Stagefright bugs*

*x509 Serialization bug*

*PingPong root – CVE-2015-3636*

rooting tool with several exploits: [https://github.com/android-rooting-tools/android\\_run\\_root\\_shell](https://github.com/android-rooting-tools/android_run_root_shell)

The first part is about the discovery of the vulnerability (CVE-2015-3636) which we leverage to achieve privilege promotion on Android devices. It is originally an access of invalid virtual address in Linux kernel found by our custom kernel syscall fuzzer (lead by @wushi of KeenTeam). And I will show how to turn it into a use-after-free bug on PING socket object in the kernel. The root cause of this bug will be revealed, which reflects certain insecurity of the Android kernel compared with the Linux kernel currently.



WEN XU

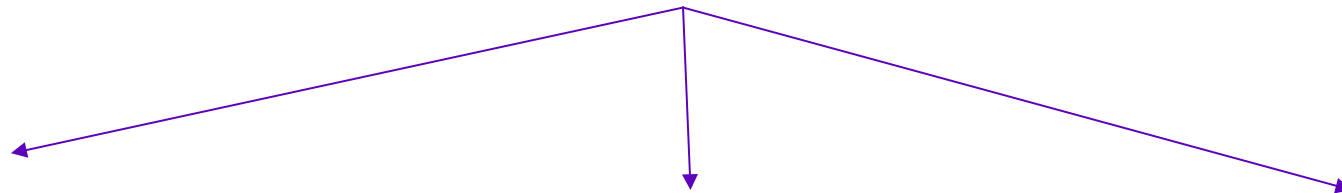
AWESOME!



# Threat Model

## BSP (Android core OS)

Kernel, nLibs, HAL, nDaemons



T1 - Root  
Privilege Escalation

T2 - Execution arbitrary  
code

T3 - Denial of Service  
(services)

# is it really bigger?



- Integer overflows are mitigated!
- Heap spray becomes a serious issue!
- New addressing @64bit makes more tricky writing realible exploits
- No Thum anymore?! Less ROP fun...
- New instruction and compiler optimizations
- At the moment, less HW availability!



## Attack surface for rooting

Nothing changed... what are our fuzzing targets?

- **Syscalls.**
- **Drivers, IOCTLs**
- All the **communications between Userspace and Kernel Space**: i.e. Motochopper is a mmap problem on the /graphics/fb0



Surface  
fuzzing

**Vulns are deep in the code!**

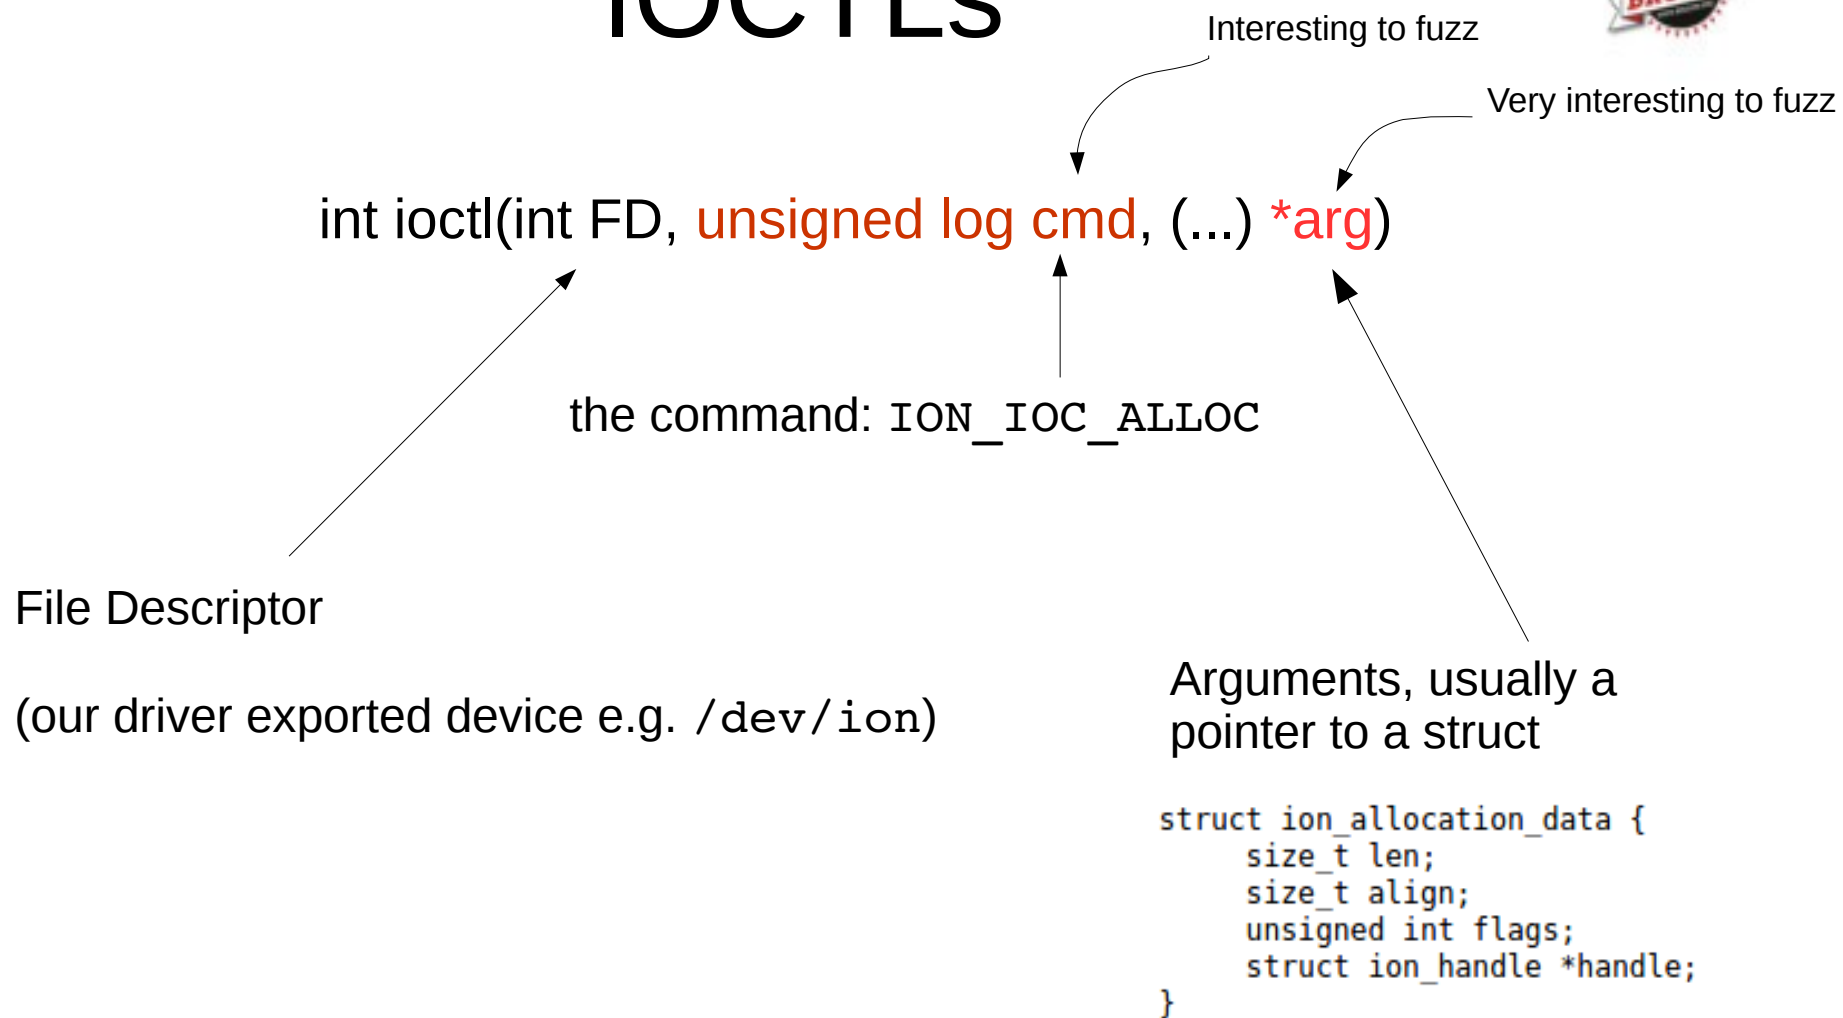


Source Code  
analysis

access\_ok — Checks if a user space pointer is valid  
get\_user — **Get a simple variable from user space.**  
put\_user — Write a simple value into user space.  
\_\_get\_user — Get a simple variable from user space, with less checking.  
\_\_put\_user — Write a simple value into user space, with less checking.  
\_\_copy\_to\_user — Copy a block of data into user space, with less checking.  
\_\_copy\_from\_user — Copy a block of data from user space, with less checking.  
strlen\_user — Get the size of a string in user space.  
\_\_strncpy\_from\_user — Copy a NUL terminated string from userspace, with less checking.  
strncpy\_from\_user — Copy a NUL terminated string from userspace.  
clear\_user — Zero a block of memory in user space.  
\_\_clear\_user — Zero a block of memory in user space, with less checking.  
strlen\_user — Get the size of a string in user space.  
copy\_to\_user — Copy a block of data into user space.  
copy\_from\_user — **Copy a block of data from user space.**



# IOCTLs



ION is a buffer sharing drivers, similar to DMABUF but with an userspace device. It replaces PMEM in Android for memory pools provisioning



# Propagation of the taint

```
31 static long my_ioctl(struct file *f, unsigned int cmd, unsigned long arg)
32 #endif
33 {
34     query_arg_t q;
35
36     switch (cmd)
37     {
38         case QUERY_GET_VARIABLES:
39             q.status = status;
40             q.dignity = dignity;
41             q.ego = ego;
42             if (copy_to_user((query_arg_t *)arg, &q, sizeof(query_arg_t)))
43             {
44                 return -EACCES;
45             }
46             break;
47         case QUERY_CLR_VARIABLES:
48             status = 0;
49             dignity = 0;
50             ego = 0;
51             break;
52         case QUERY_SET_VARIABLES:
53             if (copy_from_user(&q, (query_arg_t *)arg, sizeof(query_arg_t)))
54             {
55                 return -EACCES;
56             }
57             status = q.status;
58             dignity = q.dignity;
59             ego = q.ego;
60             break;
61         default:
62             return -EINVAL;
63     }
64
65     return 0;
66 }
67
```

- cmd is switched
- copy\_from\_user of \*arg
- Inner code to call copy\_from\_user again for inner pointers of \*arg





## Vulnerability Explained

- Userspace memory mapped with shellcode and NOPs

```
/* look at the objdump of shellcode to see the correct offset */
memcpy(map + MMAP_OFF, (unsigned char *)shellcode + 8 /* offseting to the __transgressor */, 30 *
```

- Over again, IOCTL call to binder

```
fprintf(stderr, "[!] binder local root exploit\n[!] (c) piotr szerman\n");
```

```
fd = open("/dev/binder", O_RDWR);
```

```
        ioctl(fd, BINDER_WRITE_READ, &bwr);
        close(fd);
```

- But &bwr is properly defined.

```
/* Binder transaction request format */
```

```
struct binder_write_read {
    signed long    write_size;      /* bytes to write */
    signed long    write_consumed; /* bytes consumed by driver */
    unsigned long  write_buffer;
    signed long    read_size;      /* bytes to read */
    signed long    read_consumed;  /* bytes consumed by driver */
    unsigned long  read_buffer;
} bwr;
```

```
bwr.write_size = 0;
bwr.write_consumed = 0;
bwr.read_size = 1;
bwr.read_consumed = 0;
/* targeting the aperture bet
bwr.read_buffer = (unsigned 1
```



## Review of access\_ok (updated)

- size verification, access to the memory, more checks can be done before entering sensitive areas like copy\_from\_user or copy\_to\_user
- Access\_ok is called by copy\_from\_user in new kernels and only \_\_copy\_from\_user is direct without using it
- Be carefull to which function you are using and if your access\_ok is vulnerable! Access\_ok has an important vuln in old kernels.

```
/*
 * extract the type and number bitfields, and don't decode
 * wrong cmds: return ENOTTY (inappropriate ioctl) before access_ok( )
 */
if (_IOC_TYPE(cmd) != SCULL_IOC_MAGIC) return -ENOTTY;
if (_IOC_NR(cmd) > SCULL_IOC_MAXNR) return -ENOTTY;

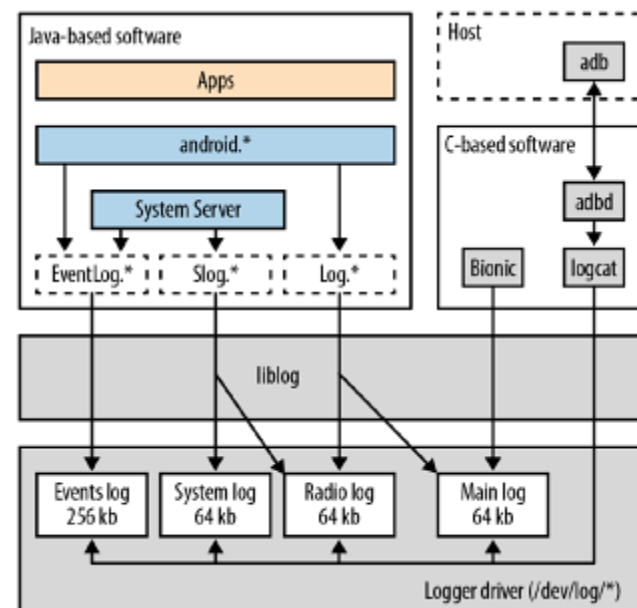
/*
 * the direction is a bitmask, and VERIFY_WRITE catches R/W
 * transfers. `Type' is user-oriented, while
 * access_ok is kernel-oriented, so the concept of "read" and
 * "write" is reversed
 */
if (_IOC_DIR(cmd) & _IOC_READ)
    err = !access_ok(VERIFY_WRITE, (void __user *)arg, _IOC_SIZE(cmd));
else if (_IOC_DIR(cmd) & _IOC_WRITE)
    err = !access_ok(VERIFY_READ, (void __user *)arg, _IOC_SIZE(cmd));
if (err) return -EFAULT;
```

# Androidization of the Kernel



- Wakelocks
- Low memory killer
- Binder
- Anonymous shared memory (ashmem)
- Alarm (extending Real Time Clock)
- Logger
- Ram Console

```
CONFIG_ASHMEM=y
CONFIG_STAGING=y
CONFIG_ANDROID=y
CONFIG_ANDROID_BINDER_IPC=y
CONFIG_ANDROID_LOGGER=y
CONFIG_ANDROID_RAM_CONSOLE=y
CONFIG_ANDROID_LOW_MEMORY_KILLER=y
```





## Research strategy

### Dumb fuzzing (dronity)

- General random syscall, random parameters, non targeted
- Targeted syscall, random parameters
- General IOCTLs against all possible devices (and netlink)
- General IOCTLs against a specific driver
- Specific IOCTLs commands against the specific driver - most effective code testing approach

**unprivileged user** (shell user via simple adb access)

**privileged user** (only in privileged mode when the driver was not world read/writable)

### Fuzzing 2.0 (working on..)

- In line kernel fuzzing with color-taint approach - using the **Zion**
- Userspace ioctls calls hooking (in LD-PRELOAD) + recursive lookup **Android Userspace fuzzer**.

### Reverse Engineering

# Reverse engineering



IDA - /root/rope

File Edit Jump Search View Debugger Options Windows Help

Local Linux debugger

Library function Data Regular function Unexplored Instruction External symbol

Functions window

- Function name
- \_init\_proc
- \_wprintf
- \_\_\_gmon\_start\_\_
- \_\_\_libc\_start\_main
- \_start
- \_\_do\_global\_dtors\_aux
- frame\_dummy
- main
- \_\_libc\_csu\_init
- \_\_libc\_csu\_fini
- \_\_i686\_get\_pc\_thunk\_bx
- \_\_do\_global\_ctors\_aux
- \_term\_proc

Graph overview

100.00% (-274,-40) (834,174) 00000404 08048404: main+10

AU: idle Down Disk: 148GB

root@Phlegethon: ~ IDA - /root/rope

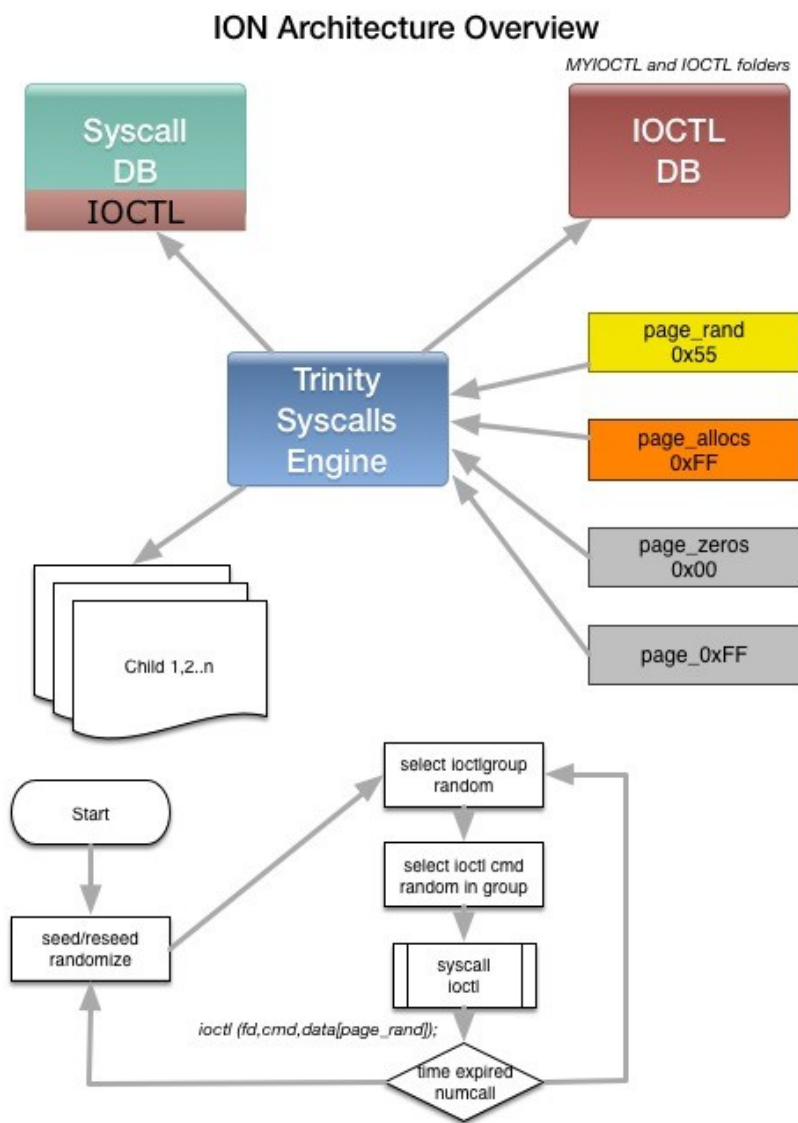
```
; Attributes: bp-based frame
public main
main proc near
push    ebp
mov     ebp, esp
and     esp, 0FFFFFFF0h
sub     esp, 10h
mov     eax, ds:debug
test    eax, eax
jz      short loc_8048412

mov     dword ptr [esp], offset unk_80484F0
call    _wprintf

loc_8048412:
mov     eax, 0
leave  eax, 0
retn
main endp
```



# Trinity arch high level view



Trinity is a Linux Kernel fuzzer developed by Dave Jones who did a splendid job compared to older projects like “iknowyou”, etc.

Based on the idea of fuzzing all the syscalls, was used in the past to discover several vulnerabilities.

Attention on it dropped since a couple of years in which the amount of vuln found dropped down.



## Other features of trinity

- Multi thread (forks on jobs)
- Enumeration of all possible system fd (not only file descriptors, so it gets almost everything)
- Extended logging now
- Extended support for kernel perfmon (excluded from dronity)
- Extended support for call\_dumps



## Trinity internals

- Definition of ARCH\_IS\_BIARCH

this parameter will enable all the 32 bit syscalls on 64 bit kernels.

32-on-32 will just use syscall() directly from do\_syscall() because do32bit flag is biarch only.

- Generic syscall is performed like this:

```
if (rec->do32bit == FALSE)
    ret = syscall(call, rec->a1, rec->a2, rec->a3, rec->a4, rec->a5, rec->a6);
else
    ret = syscall132(call, rec->a1, rec->a2, rec->a3, rec->a4, rec->a5, rec->a6);
```

- Generic initialization

- Syscalls params are generic before created as follow:

```
rec->a1 = (unsigned long) rand64();
rec->a2 = (unsigned long) rand64();
rec->a3 = (unsigned long) rand64();
rec->a4 = (unsigned long) rand64();
rec->a5 = (unsigned long) rand64();
rec->a6 = (unsigned long) rand64();
```



## Trinity internals



- Generic SANITIZE (where Trinity gets really awesome...)

generic\_sanitize look into the *syscall* definition and *argtypes* to decide what they are:  
e.g. ARG\_FD, ARG\_ADDRESS

note: ARG\_MMAP and ARG\_ADDRESS will return the cool stuff pages generated

- for *args* of *ioctls* nothing is defined, default is unsigned long rand64()!!!
- Finally, customized call on the specific driver = (see specific syscall/ioctl.c) to see how

```
if (entry->sanitise)
    entry->sanitise(rec);
```

- ioctl specific sanitize
  - a) cmd is “bit” twicked (mangle\_cmd)
  - b) arg is randomly assigned address (NULL excluded) – (mangle\_arg)



# Trinity development



**Dave Jones**

kernelslack

📍 Boston, MA

✉ [davej@codemonkey.org.uk](mailto:davej@codemonkey.org.uk)

🌐 <http://www.codemonkey.org.uk>

🕒 Joined on Nov 3, 2011

+ Contributions

📖 Re

Popular repositories

📖 trinity

<http://codemonkey.org.uk/2015/07/12/future-trinity/>

"That thing we asked you to stop working on when you worked here, any chance now you've left you'll implement these features."

Then earlier this week, came the revelation that the only people prepared to fund that kind of new feature development are pretty much the worst people.

It's a project everyone wants to take from, but no-one wants to give back to.

And that's why for the foreseeable future, I'm unlikely to make public any further feature work I do on it.  
I'm done enabling assholes.

## Introducing Dronity



### The first, real, true porting of Trinity to Android

- patched to support Bionic native under android at level 21 (arm64 and 32)
- google NDK compatible (compiles with google ndk native toolchain)
- no glibc anymore, means no static-linked distribution
- improved makefile for android to work with ndk standalone (not ndk-build)
- compiled with native .h headers and bionic, memory structures, page (PAGE\_MASK, etc) are fully compliant and compatible with android/bionic
- looked into arm64 optimization and support (need to be further improved looking to each single syscall)
- support the <constructor> macro for gcc (not C11 compatible) that prevents ioctl drivers to work with Android compilers
- support for PIE and PIC needed for Android-L to run
- gets all the benefits and feature of trinity 1.4 except perfmonitoring and backtrace (working on improving backtrace, I am up to, see changelog for more)

Dronity is pure dynamic linked aarch64 code:

```
Vitos-MacBook-Pro-5:sec bulk vito$ file dronity
dronity: ELF 64-bit LSB shared object, version 1 (SYSV), dynamically linked (uses shared libs), not stripped
Vitos-MacBook-Pro-5:sec bulk vito$
```



## Dronity first and second changelog

- fixed missing types (all around), adapting code to BIONIC
- disabling backtrace feature on child (unavailable on bionic)
- adjusted FD handling process adding missing FWM enum
- Fixing performance monitor, disabling it as it is unavailable on android (need to fix include in ia64/perfmonctl.c) \*must do: review fd-perf.c
- Removing some small unsupported features from syscalls
- Hardcode NO support for HAVE\_TERMIOS2
- Fixing network stuff and unsupported protocols (i.e. x25)
- fixing syscalls-aarch64.h for missing syscalls (1)
- Working on makefile to make it compatible with google standalone toolchain
- Extensive more changes to support Android Platform 19
  - Cope with missing functionalities of bionic
  - Removed support for fd-timerfd
  - Removed several other network protocols (incompatible and useless)
  - Patched unknown functions and removed support for cpu scheduler affinity
  - Removed swapon/off syscalls
  - ...

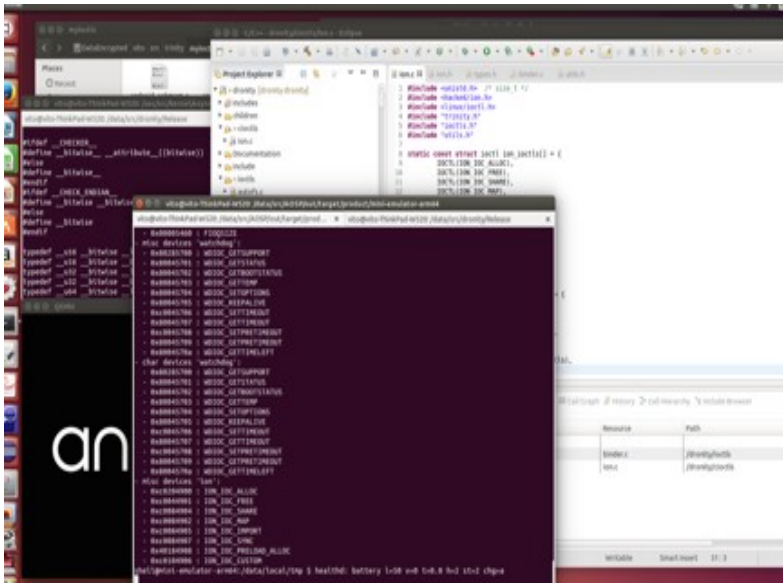


# Dronity Community Edition

Dronity is a full bionic ported version of Trinity 1.4

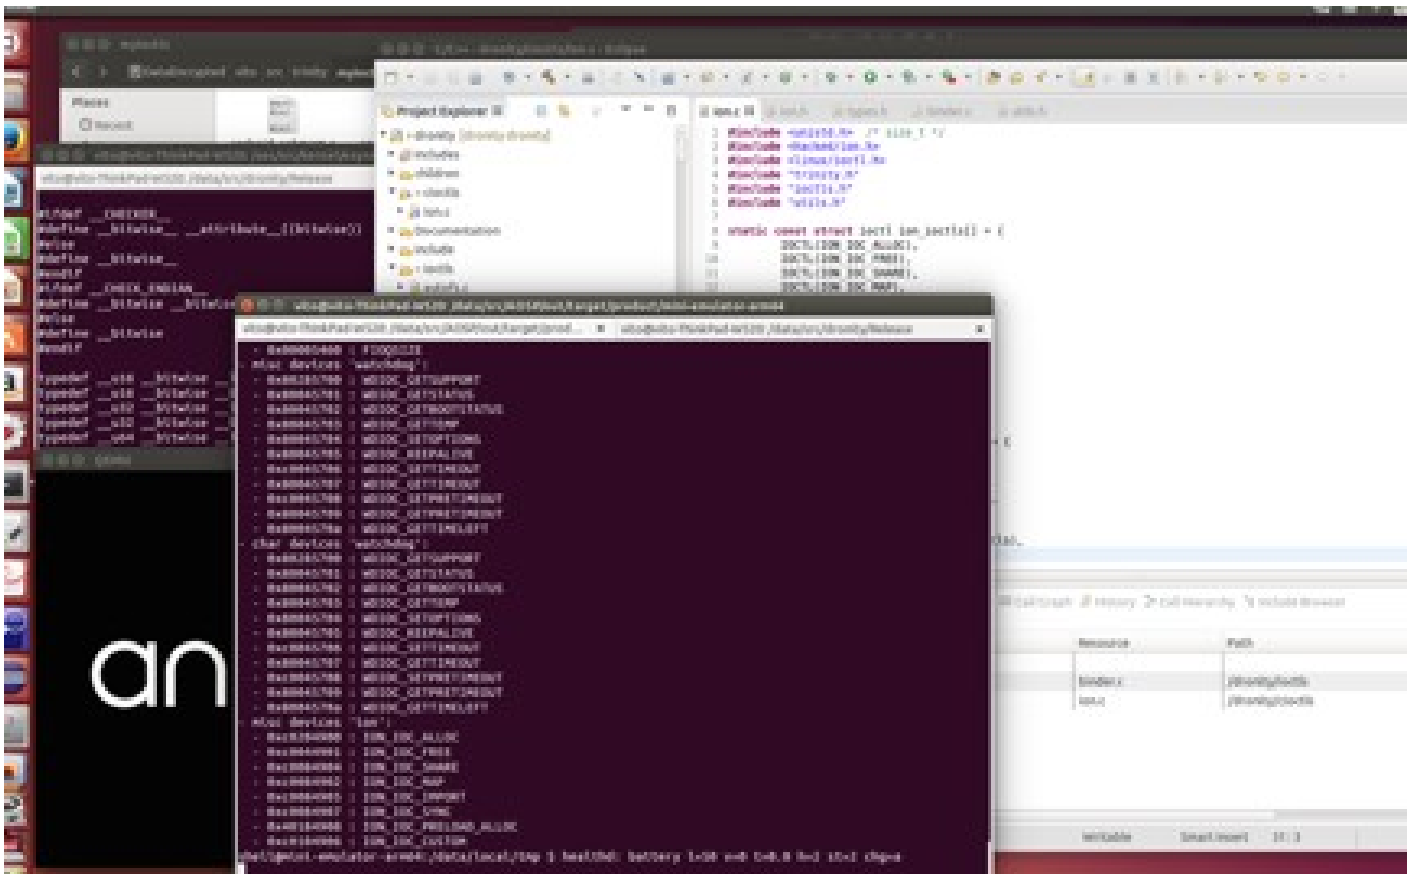
```
Vitos-MacBook-Pro-5:sec bulk vito$ file dronity
dronity: ELF 64-bit LSB shared object, version 1 (SYSV), dynamically linked (uses shared libs), not stripped
Vitos-MacBook-Pro-5:sec bulk vito$
```

- Dronity Community Edition, fancy MS style naming!
- New parameters
- Improved ioctl group management
- Improved fd management to point one file only
- Support for fixed pages
- Support for colored taint
- For this workshop: the version includes hacked kernel headers from Samsung ODDROID original code base





Develop specific IOCTL drivers to support SE IOCTLs



# Trinity new ioctl support modules



```
trinity.c  main.c  params.h  android_ion.c  fbio.c  mobicore-admin.c  »9  □
#include <unistd.h> /* size_t */
#include <linux/ion.h>
#include <linux/ioctl.h>
#include "trinity.h"
#include "ioctls.h"

static const struct ioctl ion_ioctls[] = {
    IOCTL(ION_IOC_ALLOC),
    IOCTL(ION_IOC_FREE),
    IOCTL(ION_IOC_SHARE),
    IOCTL(ION_IOC_IMPORT),
    IOCTL(ION_IOC_SYNC),
    IOCTL(ION_IOC_CUSTOM),
};

static const char *const ion_devs[] = {
    "ion",
};

static const struct ioctl_group ion_grp = {
    .name = "ion",
    .devtype = DEV_MISC,
    .devs = ion_devs,
    .devs_cnt = ARRAY_SIZE(ion_devs),
    .sanitise = pick_random_ioctl,
    .ioctls = ion_ioctls,
    .ioctls_cnt = ARRAY_SIZE(ion_ioctls),
};

REG_IOCTL_GROUP(ion_grp)
```





# Fuzzing random syscalls

```
vito@vito-ThinkPad-T440s: ~/opt/mini-emulator-arm64
vito@vito-ThinkPad-T440s: ~/opt/mini-emulator-arm64 x vito@vito-ThinkPad-T440s: ~/src/dronity64 x root@vito-ThinkPad-T440s: /opt/android-ndk-r10e/bui... x

[child0:15915] [91] mq_open(u_name=0x0, oflag=0xffffffff80800fff, mode=4107, u_attr=0xffffffffc2002030) = -1 (Bad address)
[child0:15915] [92] get_robust_list(pid=0, head_ptr=0x0, len_ptr=0xffffffff7fffffff) = -1 (Bad address)
[child0:15915] [93] mq_timedreceive(mqdes=403, u_msg_ptr=0x7f8a500000, msg_len=0x80d051, u_msg_prio=0x8000000000917e81, u_abs_timeout=0x8000000000917e85) = -1 (Bad address)
[child0:15915] [94] mq_getsetattr(mqdes=403, u_mqstat=0x0, u_o mqstat=0x8) = -1 (Bad file number)
[child0:15915] [95] getpid() = 0x3e2b
[child0:15915] [96] fsync(fd=403) = -1 (Invalid argument)
[child0:15915] [97] rt_sigtimedwait(uthese=-1, uinfo=0xfffffffffff0707, uts=0x4000000000000a2, sigsetsize=0x4000000000000a3) = -1 (Invalid argument)
[child0:15915] [98] setns(fd=405, nstype=0x44000000) = 0
[child0:15915] [99] getcwd(buf=0x7f8c0f5000, size=230) = 16
[child0:15915] [100] rt_sigtimedwait(uthese=0x8800000000000100, uinfo=0x7f88300000, uts=0x0, sigsetsize=0x8) = -1 (Bad address)
[child0:15915] [101] setgid(gid=0xfffffffffbfbffd7fb) = -1 (Operation not permitted)
[child0:15915] [102] write(fd=405, buf=0x7f8c0f5000, count=3771) = 8
[child0:15915] [103] sched_setaffinity(pid=0, len=0, user_mask_ptr=0x1) = -1 (Invalid argument)
[child0:15915] [104] kexec_load(entry=0xfffffffffffff, nr_segments=2, segments=0xffffffffffe00007, flags=0x2) = -1 (Function not implemented)
[child0:15915] kexec_load (104) returned ENOSYS, marking as inactive.
[child0:15915] [105] munlockall() = -1 (Invalid argument)
[child0:15915] [106] fallocate(fd=405, mode=0x8, offset=0, len=0x7cb3000) = -1 (Operation not supported on transport endpoint)
[child0:15915] [107] fchownat(dfd=405, filename="/sys/bus/scsi/drivers/sd", user=0xffffffffc00fff000, group=0x800000000009fdc, flag=0xffffffffffff4d4ed3) = -1 (Invalid argument)
[child0:15915] [108] lookup_dcookie(cookie64=0xffffffff7ffffff14, buf=0x1, len=16) = -1 (Operation not permitted)
[child0:15915] [109] add_key(_type=0x7f8ad00000, _description=0xfffffffffff3030, _payload=0xfffffffffff3030, plen=0x3fffffff, ringid=0xffffffff00000000) = -1 (Invalid argument)
[child0:15915] [110] fcntl(fd=405, cmd=0x406, arg=405) = 621
[child0:15915] [111] setresgid(rgid=0x9000000010000000, egid=0xeef8000, sgid=0xe3981f160c548030) = -1 (Operation not permitted)
[child0:15915] [112] rt_sigsuspend(unewset=0x800000009215f000, sigsetsize=0xfffffffffe0000fff) = -1 (Invalid argument)
[child0:15915] [113] fremovexattr(fd=403, name=0x4) = -1 (Bad address)
[child0:15915] [114] sendto(fd=403, buff=0xffffffffc00080000, len=61, flags=0xd6b1, addr=0x7f8b413000, addr_len=394) = -1 (Socket operation on non-socket)
[child0:15915] [115] timer_settime(timer_id=-205, flags=0x1, new_setting=0x1, old_setting=0x7f88d00000) = -1 (Bad address)
[child0:15915] [116] statfs(pathname="/sys/bus/serio/drivers", buf=0xa000000000000009) = -1 (Bad address)
[child0:15915] [117] unlinkat(dfd=403, pathname="/sys/bus/hid/devices", flag=0x8cbd7ad426958985) = -1 (Invalid argument)
[child0:15915] [118] sched_getparam(pid=0, param=0xffffffffc00080000) = -1 (Bad address)
[child0:15915] [119] fdasyncc(fd=403) = -1 (Invalid argument)
[child0:15915] [120] munlock(addr=0x7f8af00000, len=0xf2000) = -1 (Out of memory)
[child0:15915] [121] eventfd2(count=0, flags=0x80000) = 622
[child0:15915] [122] init_module(umod=0x0, len=255, uargs=0x8) = -1 (Operation not permitted)
[child0:15915] [123] readv(fd=403, vec=0x7f8b492e00, vlen=206) ^C[main] Bailing main loop because ctrl-c.
[watchdog] [1789] Watchdog exiting because ctrl-c.
[init] Ran 686 syscalls. Successes: 113 Failures: 573
shell@mini-emulator-arm64:/data/local/tmp $ healthd: battery l=50 v=0 t=0.0 h=2 st=2 chg=a
```





## Dronity new parameters

- f** allows to target a specific device like /dev/binder. This extend -V (no need to use -V anymore) and exclude all the other non-file related file descriptors (output, buffers, network and so on..)
- i** similar to the feature defined last year, it allows selecting a specific group of ioctl commands. this is now improved, selection of the group can happen by name or by device fd, works well with -f
- e** exclude a specific ioctl command (not yet implemented)

*to test a specific device use:*

*dronity -f /dev/binder -i binder -c ioctl*

- I** was extended to get better touch of groups and fd descriptors!
- A** specify a fixed table to send (only deadbeaf supported now for colored tainting)



# Targeting IOCTLs, improved params

```
vito@vito-ThinkPad-T440s: ~/opt/mini-emulator-arm64
vito@vito-ThinkPad-T440s: ~/opt/mini-emulator-arm64 x vito@vito-ThinkPad-T440s: ~/Desktop/deck_kerntales x root@vito-ThinkPad-T440s: /opt/android-ndk-r10e/bui... x

[child0:3293] [9340] ioctl(fd=4, cmd=0x40087705, arg=0xb600a0df) = -1 (Invalid argument)
[child0:3293] [9341] ioctl(fd=4, cmd=0x81007702, arg=0xffffffffc000080000) = -1 (Bad address)
[child0:3293] [9342] ioctl(fd=4, cmd=0x7709, arg=0x7fb6e00000) = -1 (Invalid argument)
[child0:3293] [9343] ioctl(fd=4, cmd=0x40087707, arg=0xc00090cc3f150637) = -1 (Invalid argument)
[child0:3293] [9344] ioctl(fd=4, cmd=0x40087707, arg=0xfffffe000) = -1 (Invalid argument)
[child0:3293] [9345] ioctl(fd=4, cmd=0x41007701, arg=0xffffffffc000080000) = -1 (Bad address)
[child0:3293] [9346] ioctl(fd=4, cmd=0x40047705, arg=28) = -1 (Not a typewriter)
[child0:3293] [9347] ioctl(fd=4, cmd=0x7709, arg=0xffffffffc000080000) = -1 (Invalid argument)
[child0:3293] [9348] ioctl(fd=4, cmd=0x40047703, arg=223) = -1 (Not a typewriter)
[child0:3293] [9349] ioctl(fd=4, cmd=0x770a, arg=0x9818e000) = -1 (Operation not permitted)
[child0:3293] [9350] ioctl(fd=4, cmd=0x40047703, arg=128) = -1 (Not a typewriter)
[child0:3293] [9351] ioctl(fd=4, cmd=0x7706, arg=0x9b858000) = 0
[child0:3293] [9352] ioctl(fd=4, cmd=0x40087703, arg=0xffffffffc000080000) = 0
[child0:3293] [9353] ioctl(fd=4, cmd=0x7709, arg=241) = -1 (Invalid argument)
[child0:3293] [9354] ioctl(fd=4, cmd=0x40047703, arg=0xab545008410883) = -1 (Not a typewriter)
[child0:3293] [9355] ioctl(fd=4, cmd=0x40087703, arg=0xffffffffc000080000) = 0
[child0:3293] [9356] ioctl(fd=4, cmd=0x7704, arg=0x7fb7a3f000) = 0x80000
[child0:3293] [9357] ioctl(fd=4, cmd=0x770a, arg=0xfffff000) = -1 (Operation not permitted)
[child0:3293] [9358] ioctl(fd=4, cmd=0x41007701, arg=0xffffffffc000080000) = -1 (Bad address)
[child0:3293] [9359] ioctl(fd=4, cmd=0x81007702, arg=0x9400c8) = -1 (Bad address)
[child0:3293] [9360] ioctl(fd=4, cmd=0x41007701, arg=0xfdddf3f) = -1 (Bad address)
[child0:3293] [9361] ioctl(fd=4, cmd=0x40087707, arg=0x7fb7a3f000) = -1 (Invalid argument)
[child0:3293] [9362] ioctl(fd=4, cmd=0x7709, arg=0x7fb6300000) = -1 (Invalid argument)
[child0:3293] [9363] ioctl(fd=4, cmd=0x7706, arg=0x7fb6b00000) = 0
[child0:3293] [9364] ioctl(fd=4, cmd=0x40047703, arg=0x2b3d8f) = -1 (Not a typewriter)
[child0:3293] [9365] ioctl(fd=4, cmd=0x40087705, arg=0) = 0
[child0:3293] [9366] ioctl(fd=4, cmd=0x81007702, arg=0xffffffff95959595) = -1 (Bad address)
[child0:3293] [9367] ioctl(fd=4, cmd=0x40087705, arg=0xffffffffc000080000) = -1 (Invalid argument)
[child0:3293] [9368] ioctl(fd=4, cmd=0x81007702, arg=0x7fb6700000) = 0
[child0:3293] [9369] ioctl(fd=4, cmd=0x40087708, arg=0x7fb7a3b000) = -1 (Invalid argument)
[child0:3293] [9370] ioctl(fd=4, cmd=0x40087705, arg=0xffffffffc000080000) = -1 (Invalid argument)
[child0:3293] [9371] ioctl(fd=4, cmd=0x40087707, arg=0x40000) = -1 (Invalid argument)
[child0:3293] [9372] ioctl(fd=4, cmd=0x40087708, arg=64) = -1 (Invalid argument)
[child0:3293] [9373] ioctl(fd=4, cmd=0x40047703, arg=102) = -1 (Not a typewriter)
[child0:3293] [9374] ioctl(fd=4, cmd=0x40047705, arg=0x7fffd8db) = -1 (Not a typewriter)
[child0:3293] [9375] ioctl(fd=4, cmd=0x7706, arg=0x7fb6b00000) = 0
[child0:3293] [9376] ioctl(fd=4, cmd=0x81007702, arg=0xbc39e216) = -1 (Bad address)
[child0:3293] [9377] ioctl(fd=4, cmd=0x40087705, arg=0xffffffffc000080000) = -1 (Invalid argument)
[child0:3293] [9378] ioctl(fd=4, cmd=0x81007702, arg=0x1c271000) = -1 (Bad address)
[child0:3293] [9379] ioctl(fd=4, cmd=0x40047705, arg=0xffffffff7fffffff) = -1 (Not a typewriter)
[child0:3293] [9380] ioctl(fd=4, cmd=0x40087705, arg=0x48235000) = -1 (Invalid argument)
[child0:3293] [9381] ioctl(fd=4, cmd=0x7704, arg=0xffffffffc000080000) = 0x80000
[child0:3293] [9382] ioctl(fd=4, cmd=0x40087708, arg=0xffffffffc000080000) = -1 (Invalid argument)
[child0:3293] [9383] ioctl(fd=4, cmd=0x7709, arg=0xad54000) = -1 (Invalid argument)
```



Is that everything you need?



# Debugging the kernel

- kgdb
  - Based on gdb
  - Client-server approach
  - Need to patch the UART driver (support for polling)
- P-trace
- Kernel Printk
- last\_kmsg
- Portmortem stats /data/dropbox
- Fuzzing tool log
- JTAG debugging – Segger Jlink, Trace32
  - Allows low level debugging for the bootloader



# With a bit of luck, monitoring the easy way...

- 1) serial console... almost a must!
- 2) RAM CONSOLE, `/proc/last_kmsg`
- 3) After Android 5, kernel 3.10 `> /sys/fs/pstore`

# Setting up a devel environment



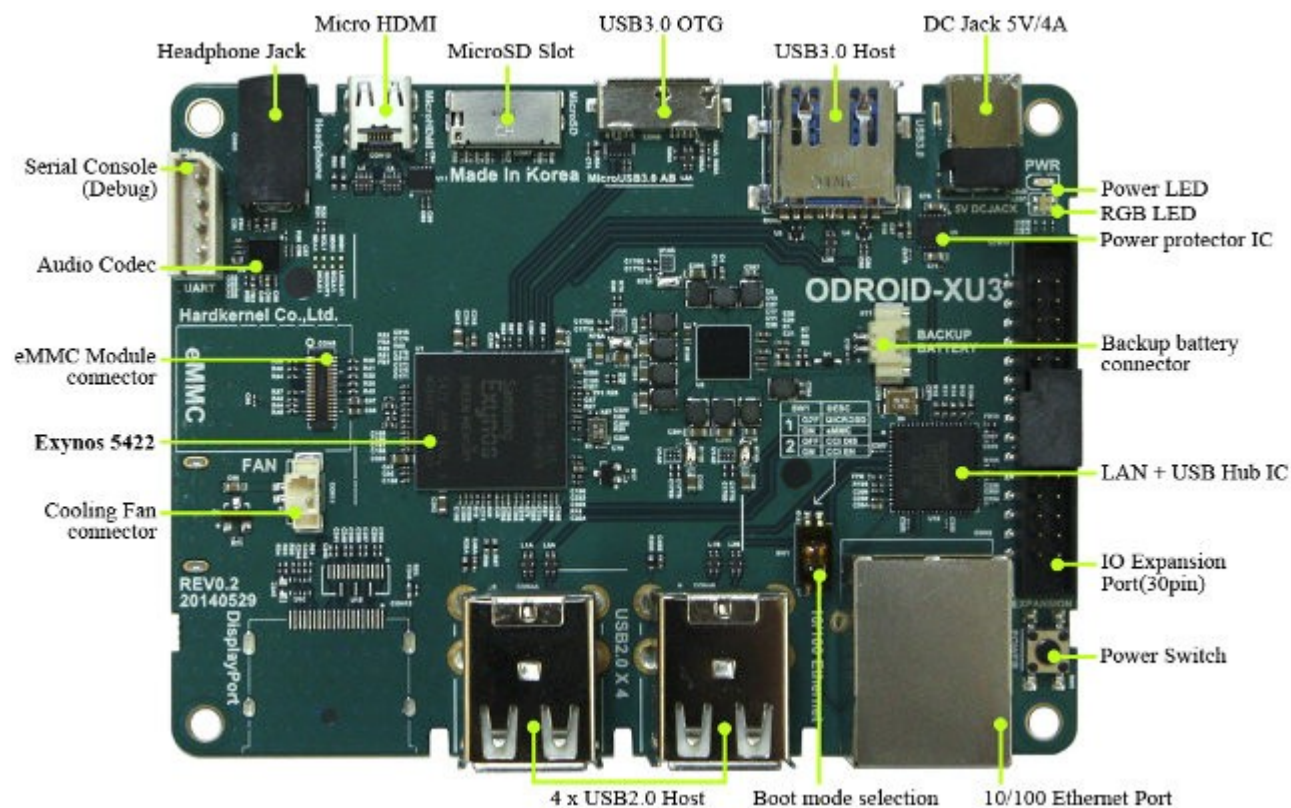
- Android SDK
- Android NDK
- Keep an eye on Linaro toolchains
- Understand about platforms and bionic
- Reading the code:
  - Eclipse
  - Understand
- Debugging: gdb.. eclipse?!

QUICK OVERVIEW DEMO

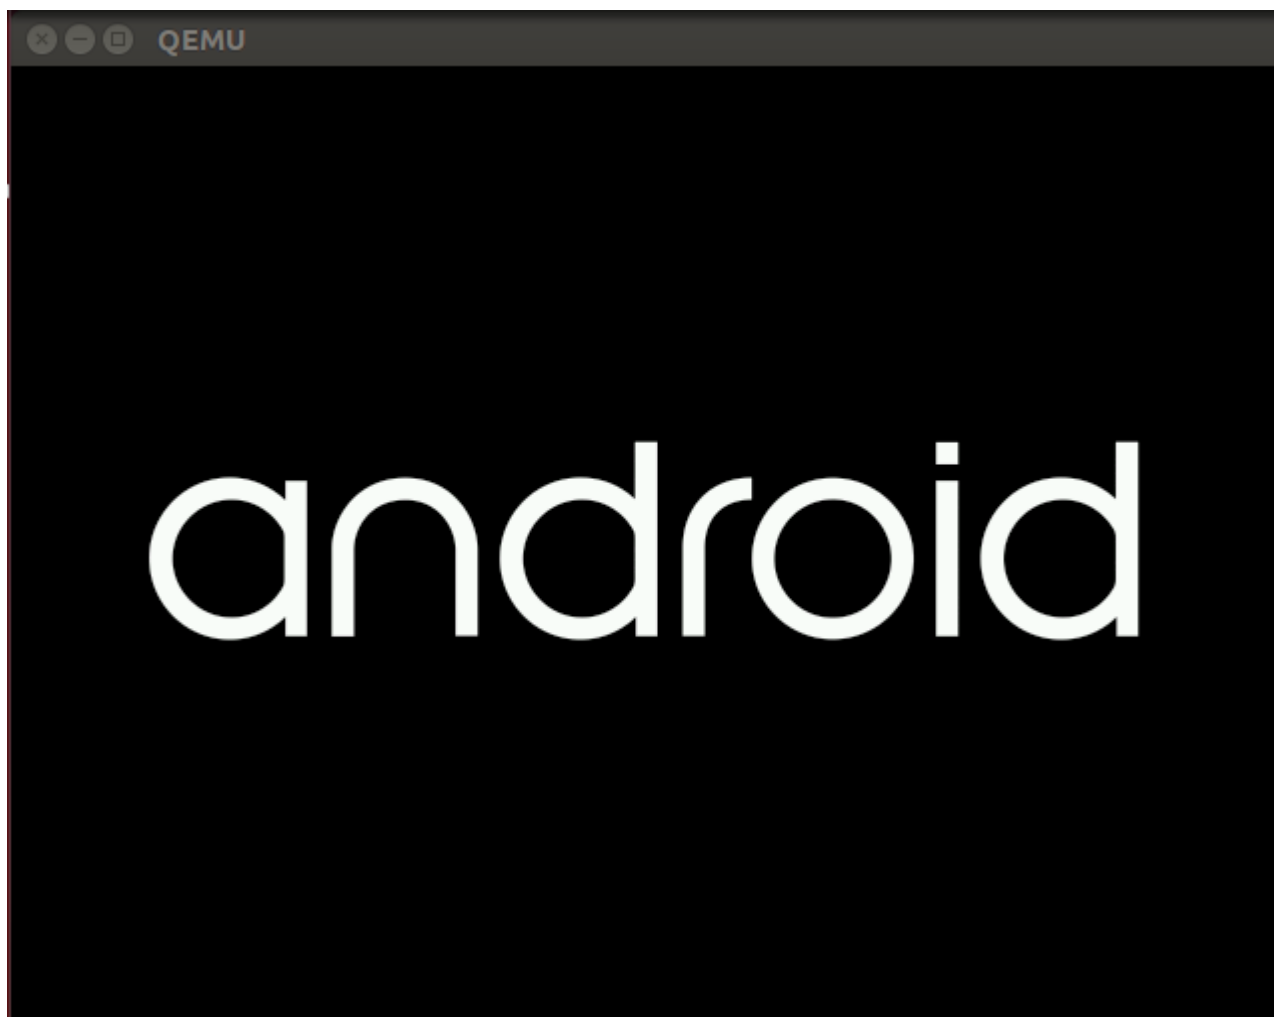




# ODROID XU3



# And for the aarch64?







# Fuzzing demo

# Compiling the kernel...



- make yourboard\_defconfig ARCH=arm64
  - Check inside your arch/<whatever>/configs what is exactly your config
  - This will generate .config file
- make CROSS\_COMPILE=your-kernel-toolchain ARCH=arm64-j8
- Flash it on the board, “fuse” to use Samsung terminology! Easy with fastboot!
- Note: you need the toolchain in the path, Linaro toolchain is great to compile kernels. You can get it from apt repos or manually from the Linaro site.

DEMO

# Cross Compile Dronity



- 1) get Android NDK from Google
- 2) run make-standalone-toolchain to get a standalone compiler

```
./make-standalone-toolchain.sh --toolchain=aarch64-linux-android-4.9 --platform=android-21 --install-dir=/opt/toolchain  
--ndk-dir=/opt/android-ndk-r10c --system=linux-x86_64
```

- 3) important: set your toolchain bin in env PATH, add it to .bashrc

```
echo "PATH=$PATH:/opt/toolchain/bin" >> ~/.bashrc  
source ~/.bashrc
```

- 4) compile dronity, you are now ready to go!

```
make CROSS_COMPILE=aarch64-linux-android- CC=aarch64-linux-android-gcc
```

(specify here the type of compiler you want, this example is based on aarch64 compiler for android version 4.9 ndk version 10c from Google)

DEMO



# Test it on the board

- `adb push dronity /data/local/tmp`
- Use the console or the adb shell
- `./dronity --help`
- `./dronity -l` (print out IOCTLs groups)

FIRST FUZZING DEMO!



# What about aarch64...

# Fuzzing aarch64 ranchu



- QEMU emulator64 is there but.. not there yet!
- There is a long story about “ranchu where are you”, article on my blog  
<http://restart-thinking.vitorallo.com/2014/11/ranchu-where-are-you-kernel-and.html>
- Need to compile emulator system images from AOSP >5
- Compiling the AOSP:  
source build/setenvironment.sh  
lunch (pick the mini\_emulator64 which does not work :) :)
- My patch still works on 5.1.1



# Ready to go?

- Go for a run, look for a girl...

```
Which would you like? [aosp_arm-eng] 14  
=====
```

PLATFORM_VERSION_CODENAME=REL
PLATFORM_VERSION=5.1.1
TARGET_PRODUCT=mini_emulator_arm64
TARGET_BUILD_VARIANT=userdebug
TARGET_BUILD_TYPE=release
TARGET_BUILD_APPS=
TARGET_ARCH=arm64
TARGET_ARCH_VARIANT=armv8-a
TARGET_CPU_VARIANT=generic
TARGET_2ND_ARCH=arm
TARGET_2ND_ARCH_VARIANT=armv7-a-neon
TARGET_2ND_CPU_VARIANT=cortex-a15
HOST_ARCH=x86_64
HOST_OS=linux
HOST_OS_EXTRA=Linux-3.16.0-49-generic-x86_64-with-Ubuntu-14.04-trusty
HOST_BUILD_TYPE=release
BUILD_ID=LMY48M
OUT_DIR=out

```
=====
```

root@vito-ThinkPad-T440s:/home/vito/src/aosp#



# Fuzzing on QEMU & Fuzzing on Nexus 9





# You got a crash, now what?

## CAUTION

This section is not recommended for VIM power users and real kernel developers

# Loading your code into Eclipse to navigate it...



It almost works, helps you navigate it... not perfect but the indexer helps and works if properly configured with compiler specs!



<http://restart-thinking.vitorallo.com/2014/11/tutorials-on-how-to-build-edit-and.html>

DEMO ACCORDING TIME



# “Understand” the code...

```
11
12 #ifndef _UAPI_LINUX_ASHMEM_H
13 #define _UAPI_LINUX_ASHMEM_H
14
15 #include <linux/ioctl.h>
16
17 #define ASHMEM_NAME_LEN 256
18
19 #define ASHMEM_NAME_DEF "dev/ashmem"
20
21 /* Return values from ASHMEM_PIN: Was the mapping purged while unpinned? */
22 #define ASHMEM_NOT_PURGED 0
23 #define ASHMEM_WAS_PURGED 1
24
25 /* Return values from ASHMEM_GET_PIN_STATUS: Is the mapping pinned? */
26 #define ASHMEM_IS_UNPINNED 0
27 #define ASHMEM_IS_PINNED 1
28
29 struct ashmem_pin {
30     __u32 offset; /* offset into region, in bytes, page-aligned */
31     __u32 len; /* length forward from offset, in bytes, page-aligned */
32 };
33
34 #define __ASHMEMIOC 0x77
35
36 #define ASHMEM_SET_NAME _IOW(__ASHMEMIOC, 1, char[ASHMEM_NAME_LEN])
37 #define ASHMEM_GET_NAME _IOR(__ASHMEMIOC, 2, char[ASHMEM_NAME_LEN])
38 #define ASHMEM_SET_SIZE _IOW(__ASHMEMIOC, 3, size_t)
39 #define ASHMEM_GET_SIZE _IO(__ASHMEMIOC, 4)
40 #define ASHMEM_SET_PROT_MASK _IOW(__ASHMEMIOC, 5, unsigned long)
41 #define ASHMEM_GET_PROT_MASK _IO(__ASHMEMIOC, 6)
42 #define ASHMEM_PIN _IOW(__ASHMEMIOC, 7, struct ashmem_pin)
43 #define ASHMEM_UNPIN _IOW(__ASHMEMIOC, 8, struct ashmem_pin)
44 #define ASHMEM_GET_PIN_STATUS _IO(__ASHMEMIOC, 9)
45 #define ASHMEM_PURGE_ALL_CACHES _IO(__ASHMEMIOC, 10)
46
47 //this is my code to declare the vuln command
48 #define ASHMEM_VULN _IOR(__ASHMEMIOC, 11, int)
49
50 /* this is the fake struct to dereference */
51 struct my_ops {
52     ssize_t (*do_it)(void);
53 };
54
55 /* this is the fake struct initialized to NULL */
56 static struct my_ops *ops = NULL;
57
58 #endif /* _UAPI_LINUX_ASHMEM_H */
59
```

Static Function ashmem\_pin\_unpin

Defined in: ashmem.c

Return Type: int

Parameters

- struct ashmem\_area \* asma
- unsigned long cmd
- void \*p

Variables

Calls

Called By

- ashmem\_ioctl

Globals Used

Macros Used

References by Flat List

Last Analysis: 7/10/15 18:51

CodeChecked: Never

# Debugging on the QEMU ranchu



- Compile the kernel with debug info:

`CONFIG_DEBUG_INFO=y`

- Run QEMU with -s param (open gdb server localhost 1234).. and that's not possible :) need to run without google wrapper:

```
/home/vito/opt/emulator64/qemu/linux-x86_64/qemu-system-aarch64 -cpu cortex-a57
-machine type=ranchu -m 2048 \
-kernel $1 -append 'console=/dev/ttyAMA0,115200' -monitor stdio -s -S \
-initrd $2/ramdisk.img -drive index=2,id=userdata,file=$2/userdata.img \
-device virtio-blk-device,drive=userdata -device virtio-blk-device,drive=cache \
-drive index=1,id=cache,file=$2/cache.img -device virtio-blk-device,drive=system \
-drive index=0,id=system,file=$2/system.img -netdev user,id=mynet \
-device virtio-net-device,netdev=mynet -show-cursor
```

Tip: console should be available by pressing Ctl+Alt+2 or by passing "-monitor stdio"

# Debugging the board over serial



- Enable KGDB on the kernel .config:

```
CONFIG_HAVE_ARCH_KGDB=y
```

```
CONFIG_KGDB=y
```

```
CONFIG_KGDB_SERIAL_CONSOLE=y
```

- Start the kernel with the option `kgdb=ttyS0,115200`

```
setenv bootargs "kgdbwait kgdboc=ttySAC2,115200"
```

- or set it later with:

```
#echo ttyS0 > /sys/module/kgdboc/parameters/kgdboc
```

- Trigger a sysrq to stop execution

```
#echo g > /proc/sysrq-trigger (or ctrl-a+f+g in minicom)
```

-



# gdb debugging on aarch64 - demo

## Vulnerable ashmem



# How to trigger the vuln

- 1) Compile trigger.c
- 2) Load it on the QEMU or Board running the vulnerable kernel
  - 1) For the board fastboot flash kernel <kernel>
  - 2) For the emulator use the launch64\_nowrap
- 3) To trigger the NULL pointer deref  
./trigger -101 (an integer < -100)



# Conclusions

- Dronity will be available on my github (only if..)
- I would really like to get cool developers input
- Try to talk more with Mr. Dave Jones
- Dreaming enough, would be nice to create a community and develop more alternative methods...
- ...